

Laboratory Earthquake Prediction

Olha Tanyuk, Daniel Davieau, Charles South and Daniel W. Engels

Southern Methodist University, Dallas TX 75205, USA

Abstract. In this paper we present a method for predicting the timing of laboratory earthquakes using machine learning. If a similar approach can be applied to improve natural earthquake prediction it will save lives. We use data collected from a laboratory experiment which exhibits similar behavior to natural earthquakes. We train a machine learning algorithm using half of the data, then predict using the other half. We compare predicted versus actual timings to measure the algorithm’s accuracy. The result shows that the timing of laboratory earthquakes can be predicted up to 16 seconds in advance with 71% accuracy. The method and result demonstrates that machine learning can help if it can be scaled from the laboratory experiment to natural earthquakes.

1 Introduction

Earthquakes cause mass destruction and loss of life. Traditional earthquake prediction methods have relied on recurrence interval based models. Because the recurrences are not constant predictions can only be made within decade spanning time windows. One such model predicted that a magnitude 6 earthquake would occur between 1985 and 1993 in the Parkfield California area but no significant event occurred until 2004 [1].

Researchers imitate natural earthquakes in the laboratory by placing rocky material between steel blocks and applying shear stress to induce slipping. Recent improvements in the instruments [2] used to measure signals have enabled the collection of larger volume data from more realistic and unpredictable laboratory earthquakes. However processing the data and detecting patterns in it has become more difficult to work with. In this paper we demonstrate that machine learning can be used to detect patterns and make predictions from realistic, unpredictable laboratory earthquake data [4].

We use data which was collected by the Los Alamos National Laboratory and provided to the public via a Kaggle competition [4]. It consists of 629 million acoustic signal observations and an accompanying record of the time remaining until a laboratory earthquake (failure) occurred [2]. We calculate additional statistical measures such as variance, kurtosis and skew for each observation. We use half of the data to train a machine learning algorithm. With the remaining half of the data, using only the acoustic signal as input we calculate a prediction of the time remaining until failure. We measure accuracy by comparing the predicted to actual remaining time to failure from the original data.

The result shows that the timing of laboratory earthquakes can be predicted up to 16 seconds in advance with 71 percent accuracy.

The data, hardware and software allows us to predict impending earthquakes. However we only know 8-16 seconds before failure. Therefore practical applications may be limited. This may prove useful but only applies to laboratory experiments. This could be used in industry perhaps in researching materials for wallboard, machine parts and others.

2 Background

In 2017 Los Alamos National Laboratory (LANL) researchers discovered a way to successfully predict Slow Slip Earthquakes (SSE) in a laboratory experiment that simulates natural conditions. The team trained a computer to pinpoint and analyze quasi-periodic seismic and acoustic signals emitted during the movements along the fault. They processed massive amounts of data and identified a particular sound pattern previously thought to be noise that precedes an earthquake. The team was able to characterize the time remaining before a laboratory earthquake at all times, using time window of 1.8 sec of the data to make each prediction, with 89% coefficient of determination [3]. This result they achieved using machine learning technique Random Forest Regression. Their results were achieved using quasi-periodic data and can not be generalized over aperiodic data.

In the lab, the team imitated a real earthquake using steel blocks interacting with rocky material (fault gouge) to induce slipping that emitted seismic sounds. An accelerometer recorded the acoustic emission emanating from the sheared layers [3]. For the first time, researchers discovered a pattern that accurately predicted when a quake would occur. The team acknowledges that the physical traits of the lab experiment (such as shear stresses and thermal properties) differ from the real world but the application of the analysis to real earthquakes to validate their results is ongoing. This method can also be applied outside of seismology to support materials' failure research in many fields such as aerospace and energy [3]. The team's lab results reveal that the fault does not fail randomly but in a predictable manner. The observations also demonstrate that the fault's critical stress state, which indicates when it might slip, can be determined using exclusively an equation of state [3]. So far seismologists and Earth scientists have mostly relied on catalogues of historical data to try to characterize the state of faults. These catalogs contain a minute fraction of seismic data, and remaining seismic data is discarded during analysis as useless noise. The authors discovered that—in the case of their laboratory faults—hidden in this noiselike data there are signals emitted by the fault that inform them of the state of the fault much more precisely than catalogues [3].

The laboratory system is a two-fault configuration that contains fault gouge material submitted to double direct shear [4]. Two fault gouge layers are sheared simultaneously while subjected to a constant normal load and a prescribed shear velocity [4]. The laboratory faults fail in repetitive cycles of stick and slip that is

meant to mimic the cycle of loading and failure on tectonic faults [4]. While the experiment is considerably simpler than a fault in Earth, it shares many physical characteristics [4].

The driving piston displaces at a very constant velocity during the inter-event time and accelerates briefly during slip [2]. An accelerometer records the acoustic emission emanating from the shearing layers [2]. The steel blocks are extremely stiff, so the deformation takes place largely in the gouge [2]. Under a broad range of load and shear velocity conditions, the apparatus stick-slips quasi-periodically for hundreds of stress cycles during a single experiment and in general follows predictions from rate and state friction [2]. The rate of impulsive precursors accelerates as failure approaches, suggesting that upcoming laboratory earthquake timing could be predicted [2]. Experimental data has 16 earthquakes. The shortest time to failure is 1.5 seconds for the first earthquake and 7 seconds for the 7th, while the longest is around 16 seconds.

3 Data

The data we analyze is provided by LANL’s 2019 Kaggle competition [4]. It was collected using a three-block assembly with two gouge layers placed in a bi-axial stress configuration. Two 5mm thick fault gouge layers were placed between the three blocks, which were held in place by a fixed normal load. The gouge material was comprised of beads with diameter 105-149 μ m. The central block was sheared at constant displacement rate. The two data streams recorded were the shear stress and the acoustic signal. While the gouge material was in a critical shear stress regime, the shear stress abruptly dropped which indicated gouge failure(a laboratory earthquake). As applied load progressively increased, the recurrence of laboratory earthquakes progressively decreased. At smaller applied loads the slips became a-periodic. The acoustic particle acceleration was measured on the central block.

LANL is certain the signal recorded for analysis is the acoustic signal emanating from the fault

Table 1. Sample of Data Provided

Index	Seismic Signal	Time to Failure
0	12	1.469099998474121
1	6	1.469099998474121
2	8	1.469099998474121
3	5	1.469099998474121
4	8	1.469099998474121

The acoustic data are integers ranging from -5515 to 5444 and have mean of 4.52. The time to failure is in seconds. We can see that acoustic data shows




Fig. 1. The magnitude of each seismic signal and the time remaining before the next laboratory earthquake. There are 16 lab earthquakes. The shortest time to failure is 1.5 seconds for the first earthquake and 7 seconds for the 7th, while the longest is around 16 seconds.

large fluctuations just before the failure and recurs cyclically. Failures can be predicted visually as cases when huge fluctuations in the signal are followed by smaller signal values. This could be useful for predicting time to failure changes from 0 to high values.

Table 2. Seismic Signal Stats

Index	Seismic Signal
count	6.29145480
mean	4.47708428
std	2.61278939
min	9.55039650
25	2.62599707
50	5.34979773
75	8.17339516
max	1.61074009




Fig. 2. The distribution of seismic signal measurements by LANL

4 Methods and Experiments

Our goal is to predict the time remaining before the next failure using only moving time windows of the acoustic data. We divide our data into windows containing 150,000 observations each (0.0375 seconds of seismic data) therefore 4194 windows. From each time window, we compute a set of 98 potentially relevant statistical features (e.g., mean, variance, kurtosis). We apply machine learning techniques such as the Random Forest Regressor, XGB Regressor, Decision Tree Regressor, LGBM Regressor, Extra Trees Regressor to the new continuous windows of values.

Similar to the LANL study we create new features separated into two main classes: Distribution of signal’s energy: we use couple of higher order moments of the acoustic data to capture the evolution of the signal’s energy. Within each time window we compute the signal’s normalized mean, minimum, maximum and higher moments (variance, skewness, kurtosis).

Precursors: the system enters a critical state when close to failure. We rely on different percentiles and thresholds to monitor this precursory activity. We use the 1st - 9th and 91th - 99th percentiles. Our thresholds measure the count of observations that the acoustic signal spends over a threshold value f_0 and under a threshold value f_1 .

In order to avoid correlation between new features we applied principal component analysis. Instead of using 98 features, we created just 10 that represented 99.9 percent of the full data variation. We use a 70/30 random split of the full time series as training and testing data sets respectively. We compute regulariza-




Fig. 3. 1% random sample from 629,143,480 observations




Fig. 4. The min value is very close to zero and the max is 16 seconds.

tion hyper-parameters for each machine learning predicting technique by random grid search based on a 3-fold cross-validation.

4.1 Data Transformations

The distribution of time to failure is right skewed. We apply a square root transformation to normalize it and improve the prediction models. It is still not ideally normal, but has improved.

4.2 Algorithms

4.3 Recursive Neural Network

LSTM is a type of RNN that helps us deal with vanishing or exploding gradients (incorrect slope). For example most of our acoustic signal observations are within x however more extreme signals occur just before failure. The extreme differences between these observations can skew the accuracy of more traditional RNN algorithms.




Fig. 5. We checked how both variables changed over time. The red line is the acoustic data and the blue one is the time to failure. On a plot above we can see, that training data has 16 earthquakes. The shortest time to failure is 1.5 seconds for the first earthquake and 7seconds for the 7th, while the longest is around 16 seconds.

LSTM separates the long term(fault slip) from short term (slow slip) Hidden state= memory from prior observations Typical RNN input+priorhiddenstate; tanh (-1,1) activation function; new hidden state LSTM is same but includes gates to determine what information is included in hidden state and what is not.

The gates are rnn's themselves.

input+priorhiddenstate; sigmoid (0,1) activation function; new hidden state. 0, 1 allows us to forget or remember Patrick Yam: LSTM cannot handle data with 150000 sequence length therefore we wavenet in the earlier layers as feature extraction and reduce the sequence length to 150.

Forget gate decides what to include 1, or exclude 0. <https://www.youtube.com/watch?v=2GNbIKTKCfE>
<https://www.youtube.com/watch?v=8HyCNIVRbSU>

4.4 Auto Regressive integrated Moving Average (ARIMA)

The data shows evidence of non-stationarity (the mean, variance change over time). We use an ARIMA model to analyze the changing means and variance. This model also can account for white noise.

ARIMA(0,0,0)

4.5 Gradient Boosting Decision Tree

5 Results

We run 5 different techniques on a train data set (70 percent of the full data) before principal component analysis and after. Principal component analysis



Fig. 6. On this zoomed-in-time plot we can see that actually the large oscillation before the failure is not quite in the last moment. There are also trains of intense oscillations preceding the large one and also some oscillations with smaller peaks after the large one. Then, after some minor oscillations, the failure occurs. Interesting thing to check is the time between high levels of seismic signal and the earthquakes. We are considering any acoustic data with absolute value greater than 1000 as a high level




Fig. 7. More than 90% of high acoustic values are around 0.31 seconds before an earthquake




Fig. 8. Distribution of Time to Failure showing right skew




Fig. 9. Distribution of Time to Failure showing improved normality after applying square root

did not improve our findings significantly, that is why we are not represent those results here. For each model we provide hyper-parameters details for future reproducibility (Table 2). When making a prediction (red curve), we emphasize that there is no past or future information considered: each prediction uses only the information within one single time window of the acoustic signal. We quantify the accuracy of our model using R2 (the coefficient of determination) and MAE (mean absolute error), applying predicting model on a 30 percent of the full data (test data).



Method	Results
<u>RandomForestRegressor</u> 'min_samples_split': 12, 'min_samples_leaf': 4, 'max_features': 'sqrt', 'max_depth': 8 <u>n_estimators=1000</u>	mean_absolute_error: 2.0322193397002044 r2_score: 0.4953230388182709 
<u>XGBRegressor</u> 'colsample_bytree': 0.9059046706163597, 'gamma': 4.395361343977617, 'learning_rate': 0.22367318773362826, 'max_depth': 21, 'min_child_weight': 69.30113404571125, 'reg_alpha': 7.681525516492065, 'subsample': 0.8919123509453631 <u>nthreads=-1,</u> <u>n_estimators=1000</u>	mean_absolute_error: 2.0634676560951557 r2_score: 0.4834777795889124 

Fig. 10. Results by Model

6 Analysis

The most accurate results with coefficient of determination 0.5 and mean absolute error 2.03 we got using Random Forest Regressor. The most important features, shown on Fig. 7, suggest to check our data for correlated features, since rolling standard deviations show the same patterns. Plus in order to improve our predicting model, we need to get rid of features that have importance equal to 0 or close to 0, such as hypermoment, minimum, maximum and some of the rolling means. Currently we are working on improving of the model to make prediction

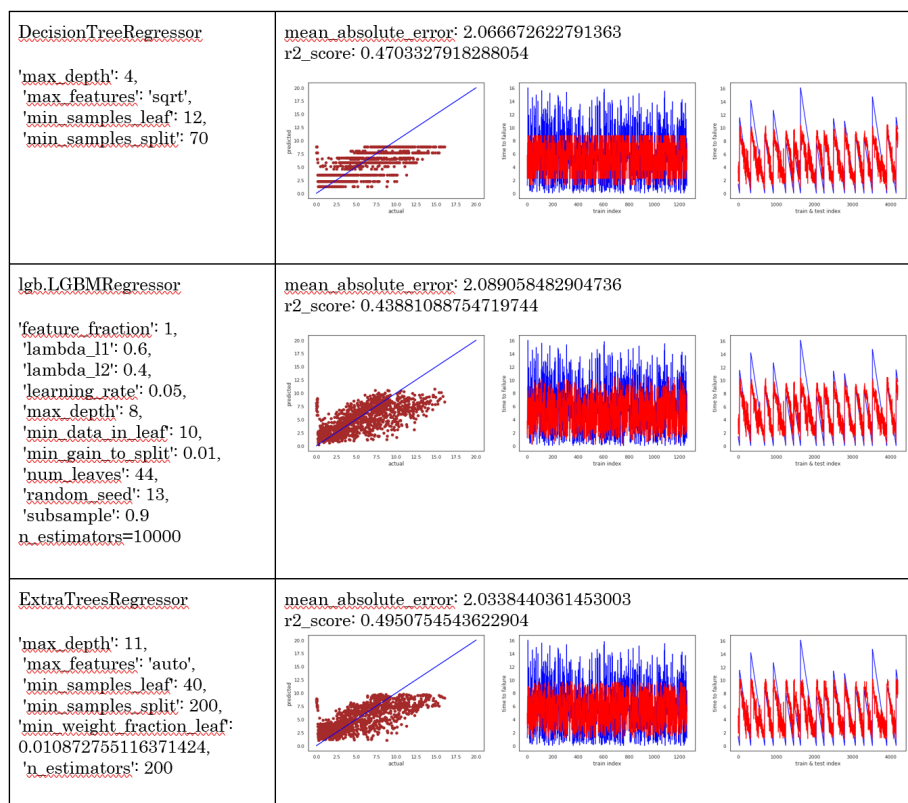


Fig. 11. Results by Model

for the beginning and the end of the one quake cycle. As you can see on a Fig. 1, seismic waves with small voltage are in our interest. Random forest model accurately predicts failure across load level, but hardly can predict outliers. It means that we still need to think about new features that we missing. Precursors that we choose during feature engineering were not good enough when the system enters a critical state close to failure.

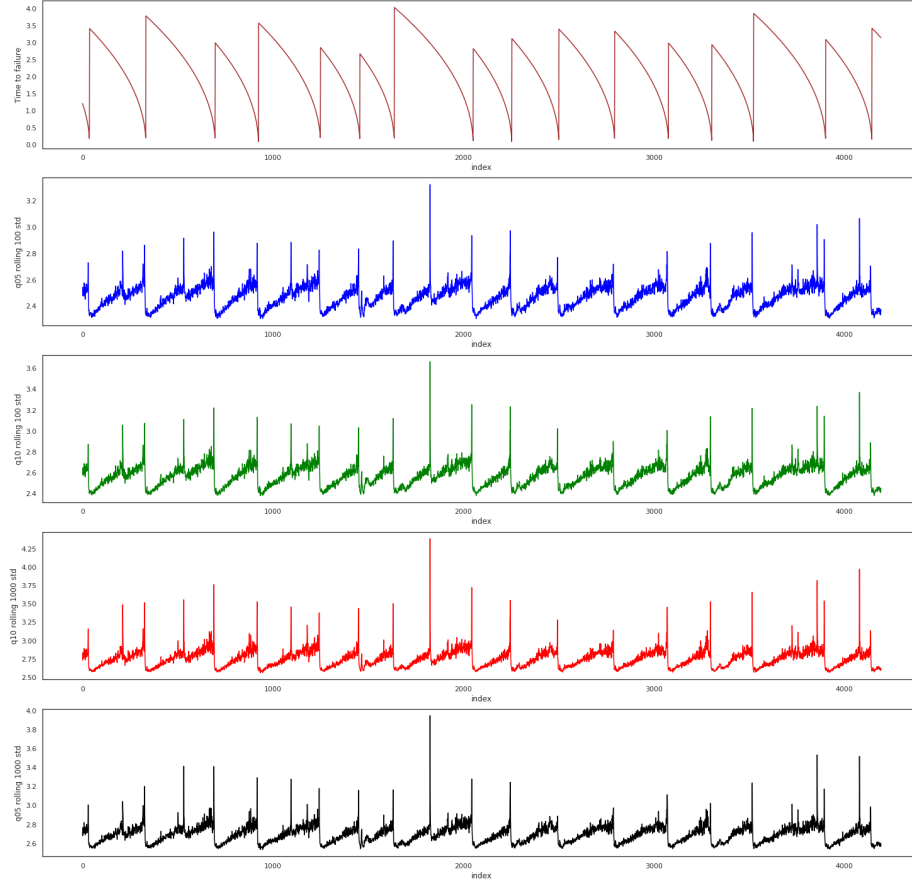


Fig. 12. Predictions by Model

7 Ethics

Our responsibility to report our results has to be weighed with our responsibility not to cause social disturbance. If the method demonstrated in this paper

is scaled and applied to predict natural earthquakes this balance must be considered. Unwarranted predictions could have affects on personal property value while failure to report warranted predictions could result in loss of life. [8].

8 Conclusions

The results show that laboratory earthquakes can be predicted with 71% accuracy up to 16 seconds in advance. The acoustic signal measurement is an indicator of imminent failure. The recurrence interval is not needed to achieve the 71% accuracy prediction.

In this study we are using only 157 seconds of data. Future work should introduce higher volumes of data to determine if accuracy can be improved. Also the combination of the acoustic signal and the recurrence interval should be tested to determine it's affect on accuracy.

References

1. Jackson, D.D.: The 2004 parkfield earthquake, the 1985 prediction, and characteristic earthquakes: Lessons for the future. *Bulletin of the Seismological Society of America* **96**(4B) (2006) 397–409 <https://doi.org/10.1785/0120050821>.
2. Rouet-Leduc, B.: Machine learning predicts laboratory earthquakes. *Geophysical Research Letters* **44**(18) (August 2017)
3. Rouet-Leduc, B.: Los Alamos machine learning discovers patterns that reveal earthquake fault behavior. Los Alamos National Laboratory, P.O. Box 1663 Los Alamos, NM 87545. *Science briefs* 2018 edn. (March 2018)
4. Kaggle: Lanl earthquake prediction. <https://www.kaggle.com/lanl> (March 2019)
5. Kious, W.J.: This dynamic earth: the story of plate tectonics. Rprt, U.S. Geological Survey, <http://pubs.er.usgs.gov/publication/7000097> (1996)
6. Ikari, M.J.: Slip weakening as a mechanism for slow earthquakes. *Nature Geoscience* **6** (May 2013) 468 <https://doi.org/10.1038/ngeo1818>.
7. Rousset, B., Bürgmann, R., Campillo, M.: Slow slip events in the roots of the san andreas fault. *Science Advances* **5**(2) (2019)
8. Ayhan Sol, H.T.: The ethics of earthquake prediction. *Science and Engineering Ethics* **10**(4) (December 2004) 655–666